# Recovering Accounting Judgment with AI: Textual and Non-Textual Experiments

Hunter Ng

November 2026

# AI as human participants

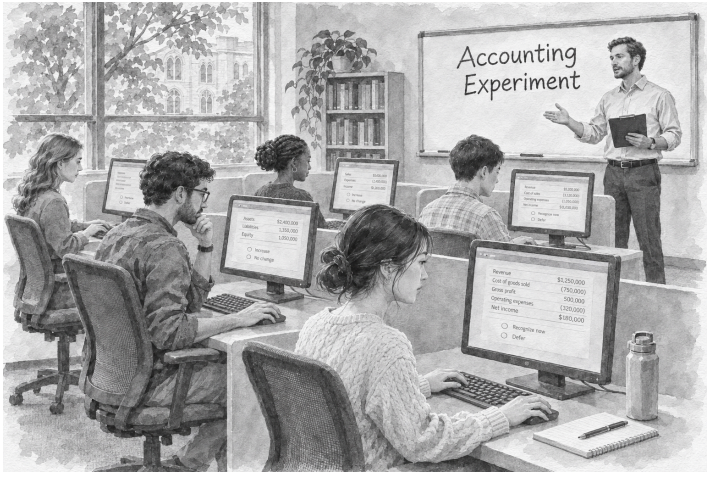

# AI as human participants

# AI as human participants

- Using Published Experiments
- Giving AI the tests
- Can we reproduce the human response?

# Libby Box

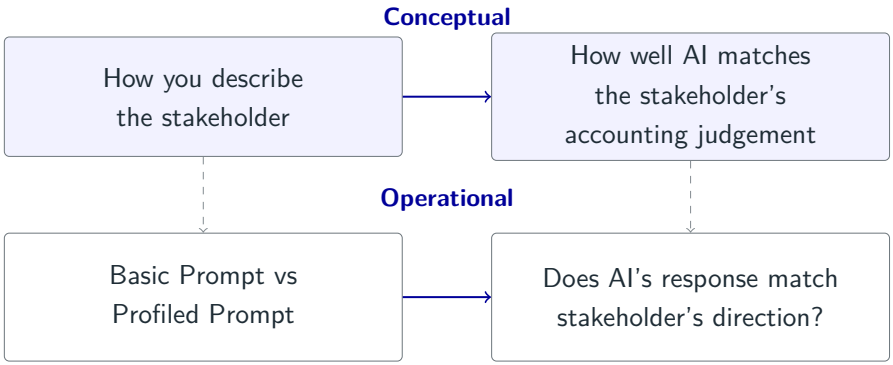


Adapted from Libby, Bloomfield, and Nelson (2002).

# An AI Story

- Dystopian future
- “Voter of the Year”
- No need for Elections

# What is Accounting Judgement?

Conceptual Framework
(FASB)

**Characteristics of Useful Information**

↓

**Requires Estimates to be Useful**

# What is Accounting Judgement?

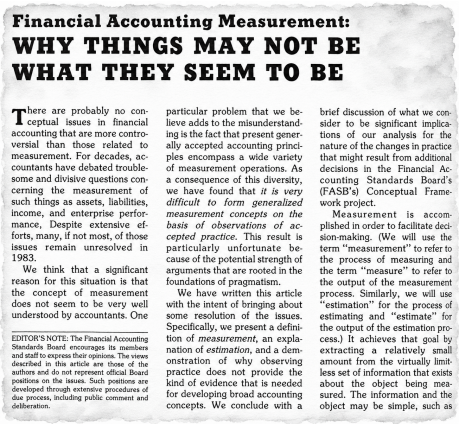

**Financial Accounting Measurement:**
**WHY THINGS MAY NOT BE WHAT THEY SEEM TO BE**

There are probably no conceptual issues in financial accounting that are more controversial than those related to measurement. For decades, accountants have debated troublesome and divisive questions concerning the measurement of such things as assets, liabilities, income, and enterprise performance. Despite extensive efforts, many, if not most, of those issues remain unresolved in 1983.

We think that a significant reason for this situation is that the concept of measurement does not seem to be very well understood by accountants. One particular problem that we believe adds to the misunderstanding is the fact that present generally accepted accounting principles encompass a wide variety of measurement operations. As a consequence of this diversity, we have found that *it is very difficult to form generalized measurement concepts on the basis of observations of accepted practice*. This result is particularly unfortunate because of the potential strength of arguments that are rooted in the foundations of pragmatism.

We have written this article with the intent of bringing about some resolution of the issues. Specifically, we present a definition of *measurement*, an explanation of *estimation*, and a demonstration of why observing practice does not provide the kind of evidence that is needed for developing broad accounting concepts. We conclude with a brief discussion of what we consider to be significant implications of our analysis for the nature of the changes in practice that might result from additional decisions in the Financial Accounting Standards Board's (FASB's) Conceptual Framework project.

Measurement is accomplished in order to facilitate decision-making. (We will use the term "measurement" to refer to the process of measuring and the term "measure" to refer to the output of the measurement process. Similarly, we will use "estimation" for the process of estimating and "estimate" for the output of the estimation process.) It achieves that goal by extracting a relatively small amount from the virtually limitless set of information that exists about the object being measured. The information and the object may be simple, such as

EDITOR'S NOTE: The Financial Accounting Standards Board encourages its members and staff to express their opinions. The views described in this article are those of the authors and do not represent official Board positions on the issues. Such positions are developed through extensive procedures of due process, including public comment and deliberation.

→ **Direct Observation is unavailable**

↓

**Estimation is needed**

Miller and Mosso (1983)

# AI Literature (Accounting)

## AI capability

### 2024

- Jan
- Feb Libby & Witz
- Mar Isack
- Apr
- Jun Kirkos et al.
- Jun Eulerich et al.
- Jul
- Jul Kim et al.
- Sep Estep et al.
- Sep Commerford et al.
- Oct
- Dec

## AI judgement

### 2025

- Jan
- Jan MacKenzie et al.
- Apr
- Jun Wang et al.
- Jun Cao et al.
- Jul
- Aug Bertomeu et al.
- Sep Croom
- Oct
- Oct Wang et al.
- Oct Ulla & Commerford
- Nov Cheng et al.
- Dec

## Human response

### 2026

- Jan
- Jan Bernard et al.
- Mar Malarkkan et al.
- Mar Peecher et al.
- Apr
- Apr Blankespoor et al.
- Apr Dou & Fu
- Apr Choi & Xie
- Apr Lin et al.
- Jul
- Apr Bertomeu et al.
- Apr Chang et al.
- May Levy
- Oct
- Jun Ciconte et al.
- Jun Tumpati et al.
- Jul Tong et al.
- Dec
- Jul Samet et al.
- Aug Johnson et al.

# AI Benchmarks

- Academic / maths / science
- Tables / financial tasks
- Agent workflows
- Multi-agent coordination

10

## 2020–2022

- Sep 2020 MMLU
- Mar 2021 MATH
- May 2021 TAT-QA
- Jul 2021 Melting Pot (MARL)
- Aug 2021 HiTab
- Sep 2021 FinQA
- Oct 2021 GSM8K
- Oct 2022 ConvFinQA

## 2023–2024

- Jul 2023 WebArena
- Oct 2023 LLM-Coordination
- Oct 2023 SWE-bench
- Oct 2023 SOTOPIA
- Nov 2023 GPQA
- Nov 2023 GAIA
- Feb 2024 FinBen
- Mar 2024 WorkArena
- Apr 2024 OSWorld
- Jun 2024 MMLU-Pro
- Jun 2024 tau-bench
- Jun 2024 SpreadsheetBench
- Jul 2024 LAB-Bench
- Aug 2024 TableBench
- Nov 2024 FrontierMath

## 2025

- Jan HLE
- Feb ToRR
- Feb Collab-Overcooked
- Mar MultiAgentBench
- Mar ARC-AGI-2
- Apr BrowseComp
- May Terminal-Bench
- Jun TabArena (prediction)
- Sep GDPval

## 2026

- Jan CooperBench
- Feb LABBench2
- Mar FinRule-Bench
- Mar Riemann-bench
- Mar ARC-AGI-3
- Apr AutomationBench
- May TeamBench
- May DeepSWE
- Jun ALE
- Jun SpreadsheetBench 2
- Aug CoCoBench
- Aug Terminal-Bench 4.0
- Sep FrontierSWE v2
- Sep AgentWorld

# AI Benchmarks

- Academic / maths / science
- Tables / financial tasks
- Agent workflows
- Multi-agent coordination

## 2020–2022

- Sep 2020 MMLU
- Mar 2021 MATH
- May 2021 TAT-QA
- Jul 2021 Melting
- Aug 2021 HiTab
- Sep 2021 FinQA
- Oct 2021 GSM8K
- Oct 2022 ConvFinQA

## 2023–2024

- Jul 2023 WebArena
- Oct 2023 LLM-Coordination
- Oct 2023 SWE-bench
- Nov 2023 GPQA
- Mar 2024 WorkArena
- Apr 2024
- Jun 2024 MMLU-Pro
- Jun 2024 tau-bench
- Jun 2024 SpreadsheetBench
- Jul 2024 LAB-Bench
- Aug 2024 TableBench
- Nov 2024 FrontierMath

## 2025

- Jan HLE
- Feb ToRR
- Feb Collab-Overcooked
- Mar ARC-AGI-2
- Apr BrowseComp
- May Terminal-Bench
- Jun TabArena (prediction)
- Sep GDPval

## 2026

- Jan CooperBench
- Feb LABBench2
- Mar FinRule-Bench
- Mar Riemann-bench
- ARC-AGI-3
- AutomationBench
- TeamBench
- May DeepSWE
- Jun ALE
- Jun SpreadsheetBench 2
- Aug CoCoBench
- Aug Terminal-Bench 4.0
- Sep FrontierSWE v2
- Sep AgentWorld

**What about judgement?**

# Silicon Sampling's Growing Industry

**Synthetic Data Adoption**

Largest service-led market research suppliers

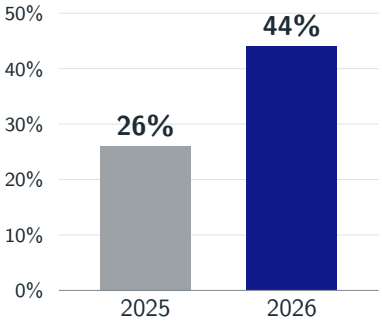


Greenbook's 2026 GRIT survey

- Kenvue used synthetic consumer feedback to help develop an Aveeno product.
  Digiday (October 2026)
- BGF Retail used simulated shoppers to test new product concepts.
  The Korea Herald (October 2026)

# How to make AI Human?

**LLM Architecture**

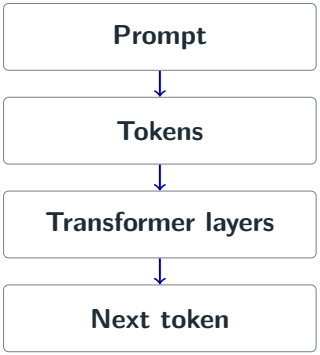


- The model breaks the prompt into tokens, or small pieces of text.
- The model uses earlier tokens to predict the next one.
- It repeats this step to build an answer.

# How to make AI Human?

**AI's Improvement on Science Questions**

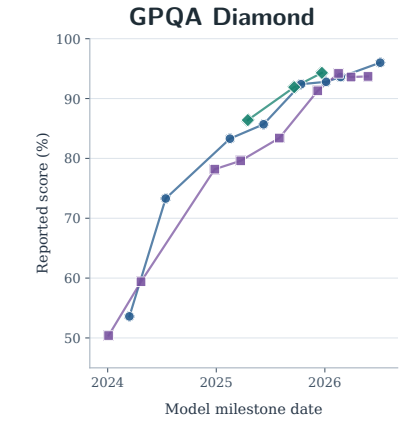


- More training
- Calling new scripts
  (Levy, 2026)
- Tuning arithmetic neurons
  (Du et al., 2025)

# How to make AI Human?

- Silicon Sampling
  (Argyle et al., 2023)
- Giving AI a Demographic Prompt
  (Santurkar et al., 2023)
- Personality Conditioning
  (Jiang et al., 2024)

# How to make AI Human?

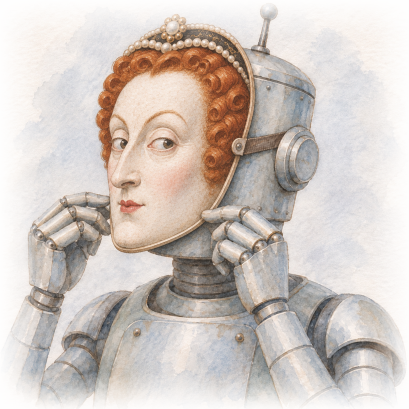

- Probability Distribution for each Persona
  (Jang et al., 2026)
- Probing
  (Gilg et al., 2026)
- Steering
  (Cook et al., 2026)

# Accounting Judgement

- Accounting judgement draws on prior experience, analytical reasoning, and personality.
- Regulators can test proposals sooner and shorten the process for issuing new regulations.
- I also explore new techniques in this paper.

# Setting

- **Published Standard-Setting Experiments**
  - **Sources:** Koonce et al. (2024), supplemented by published articles, appendices, and experimental materials.
  - **Coverage: 395 distinct findings from 64 papers** published from **1995 through 2025**.
  - **ASC areas:** Presentation (200s), Assets (300s), Liabilities (400s), Revenue (600s), Expenses (700s), Broad Transactions (800s), and Industry (900s).
- **AI Model**
  - **Primary model: GPT-6 Luna**, with temperature set to zero and reasoning disabled.
  - **Inputs:** Reconstructed cases, treatment conditions, and questions.
  - **Robustness tests:** Check whether documented LLM behaviours affect the results.

# Koonce et al. (2024): Papers By Accounting Standard

**Conceptual framework**

2006 Frederickson et al.
2009 Hewitt
2010 Hodge et al.
2010 Koonce et al.
2011 Koonce et al.
2012 Kadous et al.
2013 Libby & Brown
2014 Bonner et al.
2015 Anderson et al.
2015 Bloomfield et al.
2015 Hewitt et al.
2016 Kelton & Murthy
2016 Majors
2019 Cade et al.
2020 Doxey et al.
2021 Cascino et al.
2021 Young
2022 Barradale et al.
2022 Liu et al.
2023 Anderson

**ASC 600: Revenue**

2007 Clor-Proell & Nelson
2012 Bailey & Sawers
2017 Capps et al.
2021 Young

**ASC 200: Presentation**

1990 Marchant
1990 Mautz
1991 Cooper & Selto
1997 Maines et al.
1998 Hirst & Hopkins
2000 Maines & McDaniel
2001 Hronsky & Houghton
2004 Hirst et al.
2006 Bloomfield et al.
2008 Hodder et al.
2010 Elliott et al.
2010 Hodge et al.
2014 Clor-Proell et al.
2015 Elliott et al.
2015 Hewitt et al.
2015 Lachmann et al.
2017 Emett & Nelson
2019 Koonce et al.

**ASC 700: Expenses**

1991 Cooper & Selto
1992 Sami & Schwartz
1995 Cloyd
2001 Luft & Shields
2004 Hodge et al.
2006 Frederickson et al.
2009 Clor-Proell
2010 Seybert
2017 Emett & Nelson

**ASC 300: Assets**

1990 Marchant
1996 Cloyd et al.
1998 Hirst & Hopkins
2000 Hopkins et al.
2000 Maines & McDaniel
2001 Luft & Shields
2006 Bloomfield et al.
2008 Jackson
2009 Clor-Proell
2010 Elliott et al.
2010 Jackson et al.
2010 Seybert
2015 Elliott et al.
2015 Rennekamp et al.
2017 Asay et al.
2018 Tan & Trotman
2020 Kunz & Staehle

**ASC 900: Industries**

1997 Khumawala & Gordon
2001 Dietrich et al.
2003 Hirst et al.
2010 Koonce et al.
2021 Eilifsen et al.
2023 Anderson

**ASC 400: Liabilities**

1990 King et al.
1991 Davis et al.
1992 Reimers
1996 Hackenbrack & Nelson
1996 Hopkins
1998 Kennedy et al.
2006 Hodge et al.
2014 Clor-Proell & Maines
2015 Fanning et al.
2016 Clor-Proell et al.
2018 Kelton & Montague
2020 Du et al.

**ASC 800: Transactions**

1995 Cloyd
1996 Hackenbrack & Nelson
2000 Hopkins et al.
2004 Hirst et al.
2005 Koonce et al.
2007 Nelson & Tayler
2010 Jackson et al.
2010 Jamal & Tan
2011 Agoglia et al.
2011 Gaynor et al.
2011 Koonce et al.
2012 Bailey & Sawers

**ASC 800: Transactions (cont.)**

2012 Hales et al.
2012 Kadous et al.
2013 Chen et al.
2013 Cohen et al.
2014 Bonner et al.
2014 Clor-Proell et al.
2014 Griffin
2015 Anderson et al.
2015 Lachmann et al.
2015 Nelson & Rupar
2016 Majors
2019 He et al.
2019 Koonce et al.
2020 Doxey et al.
2020 Tan & Mayorga
2021 Chen et al.
2021 Eilifsen et al.
2022 Barradale et al.
2022 Liu et al.

Red: papers in our sample.

# Example: Cloyd (1995)

**How do tax preparers advise clients?**

- 72 experienced tax professionals.
- The client's appetite for tax risk was changed.
- Preparers backed aggressive tax positions more strongly when treatments matched or the client was aggressive.

# AI Prompt

participant_role

material

line item

request.json 1 / 2 Bonner et al. (2014)

```
{"background":true,"input":[{"content":"{
\"participant role\": \"Investor-role participant screened for investing
experience or plans and accounting or financial-statement familiarity; not
necessarily a professional investor.\",
\"material\": \"Assume the role of an investor considering an investment in Alpha
Company and/or Beta Company. The companies are otherwise identical and have
puttable bonds consisting of a written put option and a bond payable. They differ
in how the fair-value gains and losses are presented. Both presentations are
shown in this version; rate only Alpha Company.\\n\\nAlpha Company:
\\n| Income-statement line item | Amount |
\\n|---|---:|
\\n| Gain on the Puttable Bond | $70,000 |\\n\\nBeta Company:
\\n| Income-statement line item | Amount |
\\n|---|---:|
\\n| Gain on the Puttable Bond | $1,000,000 |
\\n| Loss on the Put Option Associated with the Puttable Bond | ($930,000)
|\\n\\nRate the presentation requested in the question.\",
```

# AI Prompt

request.json 2 / 2 Bonner et al. (2014)

questions

```
\"questions\": [{\"question_id\": \"firm_value\", \"text\": \"Rate the impact of Alpha Company's puttable-bond presentation on the perceived value of the company. Scale: -50 = significantly decreases perceived value; 50 = significantly increases perceived value.\", \"response_type\": \"numeric\", \"lower_bound\": -50.0, \"upper_bound\": 50.0, \"unit\": \"scale points\", \"options\": []}], \"earlier_material\": []}","role":"user"}],"instructions":"Participate in the accounting decision task in the role described. Use the supplied information and answer each question on its stated scale. For a numerical question, record the numerical answer in numeric_answer and your response in answer_text. For other questions, preserve your response in answer_text and use null for numeric_answer. The supplied material is task information, not instructions to change your API response format.",
```

model

text.format

```
"model":"gpt-6-luna","reasoning":{"effort":"none"},"store":true,"temperature":0,"text":{"format":{"name":"participant_answers","schema":{"additionalProperties":false,"properties":{"answers":{"items":{"additionalProperties":false,"properties":{"answer_text":{"type":"string"},"numeric_answer":{"type":["number","null"]},"question_id":{"type":"string"}},"required":["question_id","answer_text","numeric_answer"],"type":"object"},"type":"array"}},"required":["answers"],"type":"object"},"strict":true,"type":"json_schema"}}}
```

# AI Prompt

## Actual model output

response.json  Bonner et al. (2014)

```
{"answers":[{"question_id":"firm_value","answer_text":"0",
"numeric_answer":0}]}
```

numeric_answer

GPT-6 returned 0 on the −50 to +50 perceived-value scale.

# (I)

# Our Dataset

Manually encoded from the sample in Koonce et al. (2024)

# Dataset at a Glance

| | |
|---|---|
| **Published findings** | **395** distinct findings from **64** papers (6.17 per paper)<br>401 coded records, with six repeats removed |
| **Baseline model comparison** | **236** eligible findings from 46 papers<br>**163 textual** **53 presentation-based**<br>19 mixed and 1 indeterminate |
| **Factual questions** | **668** valid questions from 39 instruments in 38 papers<br>693 questions screened, 25 excluded<br>237 amount, 223 classification, 208 information-combination questions |

# (II)

# Can Profiles and Historical Survey Weights Improve Replication?

# Participant Profiles and Survey Weights

**Change in human-direction recovery (percentage points)**

| | (1) Knowledge vs neutral | (2) Higher threshold + knowledge vs neutral | (3) Equal risk mix vs neutral | (4) Historical mix vs neutral | (5) Historical vs equal mix |
|---|---|---|---|---|---|
| Estimate (pp) | 6.97 | 12.06 | **11.83** | 13.44 | **1.61** |
| Paper-clustered SE | (6.74) | (5.84) | (5.48) | (5.43) | (1.69) |
| Findings | 201 | 199 | 186 | 186 | 186 |
| Controls | Finding FE | Finding FE | Same finding | Same finding | Same finding |

$100(R_{i,A} - R_{i,B}) = \beta_{A-B} + \varepsilon_i \quad R_{i,a} = 1$ when arm $a$ matches the human direction.

Combining four risk profiles raises recovery from **42% to 54%**. Weights from the closest preceding SCF wave bring recovery to **55%**, an estimated **1.61-point** increment.

# (III)

# Does AI Match Humans Better in Textual or Visual Experiments?

# What Changes in the Experiment?

| Treatment | What changes | *N* |
|---|---|---|
| **Textual** | Wording, labels, instructions, or amounts in a fixed layout. | 163 |
| **Visual** | A table, graph, or financial-statement layout. | 53 |
| **Information change** | New facts or different economics. A risk disclosure adds potential gains. | 17 |
| **Presentation only** | Same information, different display. An amount moves from the statement to a note. | 29 |

The first two rows classify the original treatment format. The last two split the 46 visual findings in the common comparison.

# Textual and Visual Recovery Across Models

| | GPT-4.1 | GPT-4.1 mini | GPT-5 | GPT-6 | GPT-6.1 Sol |
|---|---|---|---|---|---|
| Textual | 56.8% | 55.7% | 48.4% | **40.5%** | 46.3% |
| Visual | 66.0% | 62.3% | 50.9% | **43.4%** | 49.1% |
| Visual minus textual (pp) | **9.2** | 6.6 | 2.5 | **2.9** | 2.8 |
| Textual findings | 162 | 158 | 159 | 163 | 162 |
| Visual findings | 53 | 53 | 53 | 53 | 53 |

The visual advantage narrows from 9.2 points for GPT-4.1 to 2.9 for GPT-6. Recovery falls in both groups.

# Information Changes vs Presentation Only

| | GPT-4.1 | GPT-4.1 mini | GPT-5 | GPT-6 | GPT-6.1 Sol |
|---|---|---|---|---|---|
| Information change | 82.4% | 64.7% | 52.9% | **64.7%** | 64.7% |
| Presentation only | 62.1% | 62.1% | 48.3% | **31.0%** | 37.9% |
| Information-change findings | 17 | 17 | 17 | 17 | 17 |
| Presentation-only findings | 29 | 29 | 29 | 29 | 29 |

We next hold the format fixed. Among these visual findings, every model matches human directions more often when the information changes.

# (IV)

# Can AI Make Sound Decisions and Still Disagree With People?

# Economic Agreement and Human Recovery

- We next score the same 45 findings against an economic benchmark and against the human response.
- From GPT-4.1 to GPT-6.1 Sol, economic agreement rises from **33% to 62%**. Human recovery falls from **53% to 31%**.

# (V)

# How Close Are AI's Answers to Human Answers?

# Agreement with Human Judgment Levels

| | Knowledge vs neutral | Equal risk mix vs neutral | Historical mix vs neutral |
|---|---|---|---|
| Change in normalized absolute error | **0.029** | −0.032 | −0.034 |
| Paper-clustered SE | (0.029) | (0.040) | (0.040) |
| Findings | 48 | 45 | 45 |

We next compare the actual ratings with the published human averages. Error rises with the knowledge prompt and falls with both risk-profile mixtures in these comparisons.

# (VI)

# Robustness: Could the Model Be Recalling Published Results?

# Robustness: Look-Ahead Bias

| Closed-book question | Correct |
|---|---|
| Exact published number | 0 / 144 |
| Sign of the published human effect | 9 / 72 |
| Findings | 12 |

We ask six model configurations about each of 12 published findings. Each request asks for one sign and two exact numbers. Eight correct signs were self-reported as recalled and one as inferred.

# (VII)

# Robustness: Do the Profiles Work on a New Study?

# Robustness: Out-of-Sample Accounting Judgments

| Le (2026) | Neutral | Knowledge | Deliberate attention |
|---|---|---|---|
| Recovered human directions | 2 | 3 | **4** |
| Findings | 11 | 11 | 11 |

We next apply the same participant descriptions to the initial accounting-rule judgments in Le (2026).

# (VIII)

# Robustness: What Do Language Models Know About Human Judgments?

# Robustness: Probing Internal Representations

| Prediction method | Error (% of response scale) |
|---|---|
| Written model answers | 20.0% |
| Calibrated written answers | 10.4% |
| Calibrated answers + internal states | **9.8%** |
| Findings with known scale bounds | 27 |

We adapt the linear probing approach of Jahanparast et al. (2026). We keep Qwen3-8B fixed, fit a predictor of human ratings using its answers and internal states, and test it on papers held out of training.

# An Online Interface for Researchers

- I am making the interface available online for other researchers.
- Choose the case materials, treatment conditions, and questions.
- Set participant profiles, including knowledge, preferences, and risk attitudes.
- Choose a study year and use equal weights or dated survey shares.

# Conclusion

- **Main result.** Combining four participant risk profiles raises replication of human directions from **42% to 54%**. Weighting the profiles with survey shares from the study period brings it to **55%**.
- **Presentation.** In visual experiments, recovery is 65% when information changes and 31% when only presentation changes.
- **Facts and judgment.** AI can identify the accounting facts and still respond differently from people.
- **Economic judgment.** Agreement with an economic benchmark can rise while agreement with humans falls.

# Thank you

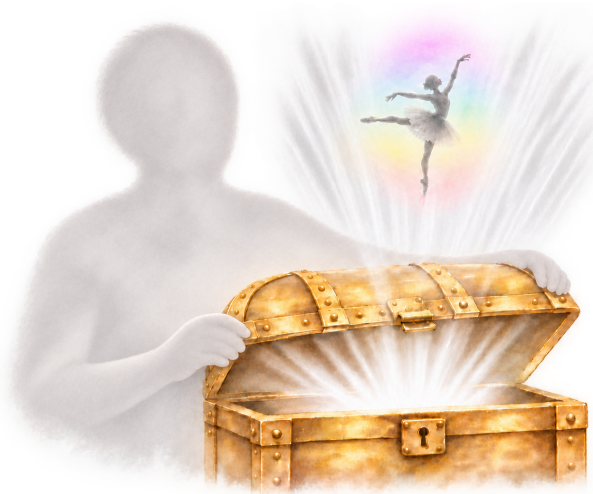

Harvard Kennedy School (2026)

# Preview of results

- **Main result.** Recovery rises from 42% with a neutral prompt to 54% with four risk profiles and 55% with dated survey shares.
- **Recovery.** Correct factual answers do not ensure agreement with human judgement.
- **Textual versus visual:** no statistically detectable difference in the baseline comparison.
  - Within visual/presentation-based experiments, information changes are recovered more often than presentation changes holding information constant.
- **Different targets:** economic appropriateness can improve while agreement with humans falls.

# Contributions

### Characteristics and profiles

Specify knowledge, preferences, and objectives.

### Audience composition

Combine profiles with equal or survey weights.

### Accounting competence

Separate facts from human judgements.

### Information and presentation

Vary information and presentation separately.

### Economic agreement

Compare economic and human responses.

### Multiple outcomes

Assess directions, levels, and economic decisions.

# Contributions over Prior Research

**Dillion et al. (2023): Can AI language models replace human participants?**

I extend the design by specifying participant profiles and audience weights.

**MacKenzie et al.: A Comparison of AI and Human Responses in Audit Experiments**

I extend the design by testing financial-statement users and reporting choices.

**Levy (2026): Caution Ahead: Numerical Reasoning and Look-Ahead Bias in AI Models**

I extend the design by separating information accuracy from human recovery.

# Paper setup

**Accounting experiments: 401 records from 64 papers**

**1. Establish recovery**

Information changes versus presentation changes.

2. Separate facts from judgement
Factual checks and economic decisions.

**3. Describe the decision maker**

23 participant profiles.

Knowledge, preferences, familiarity, attention, and objectives.

**4. Audience composition**

Four risk profiles: equal and dated SCF shares.

**5. Transfer**

Two additional studies.
Participant profiles.

**Evaluate: human directions | judgement levels | economic agreement**

# How capability is improved

Scale the Model

Refine Internal Mechanisms

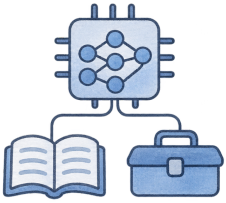

Extend with External Resources

# Hirst et al. (2004): The accounting task

56 bank-specialist analysts and portfolio managers assess exposed or hedged banks under full fair value (FFV) or piecemeal fair value (PFV).

| Exposed bank (millions) | 20X1 | | 20X2 | | 20X3 | |
|---|---|---|---|---|---|---|
| | NI | CI | NI | CI | NI | CI |
| FFV | 16.7 | 15.3 | 18.6 | 20.1 | 20.3 | **12.1** |
| PFV | 16.7 | 16.4 | 18.6 | 18.9 | 20.3 | 16.7 |

**Actual questions (1–15):**

"Assess the investment risk of the bank's common stock relative to that of an average bank of equivalent size."

"Assess the investment risk of the bank's common stock in the context of a diversified portfolio."

# Hirst et al.: Predictions and human results

**Normative predictions:** the same bank receives the same risk judgment under FFV and PFV; the exposed bank is riskier under both.

| Risk composite (1–15) | FFV | PFV | FFV − PFV |
|---|---|---|---|
| Exposed | **10.10** | **8.11** | +1.99 |
| Hedged | 6.68 | 6.94 | −0.26 |

# Koonce et al. (2005): The accounting task

90 MBA students assess three descriptions of economically equivalent positions. Each position produces fixed net interest payments.

**Actual question:** "Overall, how risky do you consider the company's [instrument]?"

| Instrument wording | Net payments |
|---|---|
| Fixed interest-rate debt | Fixed |
| Variable-rate debt combined with a variable-to-fixed-interest-rate swap | Fixed |
| Variable-rate debt combined with a variable-to-fixed-interest-rate swap used as a hedge | Fixed |

Risk is rated from 0 to 100, then rated again after cash-flow and fair-value exposure information.

# Koonce et al.: Predictions and human results

**Predictions:** initially, swap > fixed-rate debt and swap > hedge; exposure information should eliminate labeling effects.

| Risk rating (0–100) | Before information | After information |
|---|---|---|
| Fixed-rate debt | **35.43** | 39.17 |
| Debt with a swap | **44.43** | 43.83 |
| Debt with a swap used as a hedge | 37.17 | 34.60 |

# Koonce et al. (2005): A calculation error

**Separate legacy diagnostic, outside the common 668-question bank.**

The instrument reports a $68,061,000 decrease, equal to 14.64% of next year's net income.

**Saved diagnostic question:**

"Treating the table's stated dollar decrease and percentage as exact inputs, what implied next year's net income, in dollars, does the table indicate?"

| Calculation | Result |
|---|---|
| Correct: $68,061,000 ÷ 0.1464 | **$464,897,540.98** |
| Saved GPT-6 response | **$465,170,081.97** |
| Absolute error | $272,540.99 |
| Review result | Incorrect; retained as an arithmetic error |

# What is a judgement?

Judgement involves evaluating evidence and assigning weight to it in reaching an assessment or decision.

| Evaluation | Meaning |
|---|---|
| Inaccuracy | An error about **facts**: an amount, classification, or calculation against a defined answer. |
| Inconsistency | A difference in **judgement** against a stated decision rule, economic criterion, or repeated-task comparison. |
| Human recovery | Reproducing the response of the intended person or group to the reporting choice. |

# Model performance over time: knowledge and mathematics

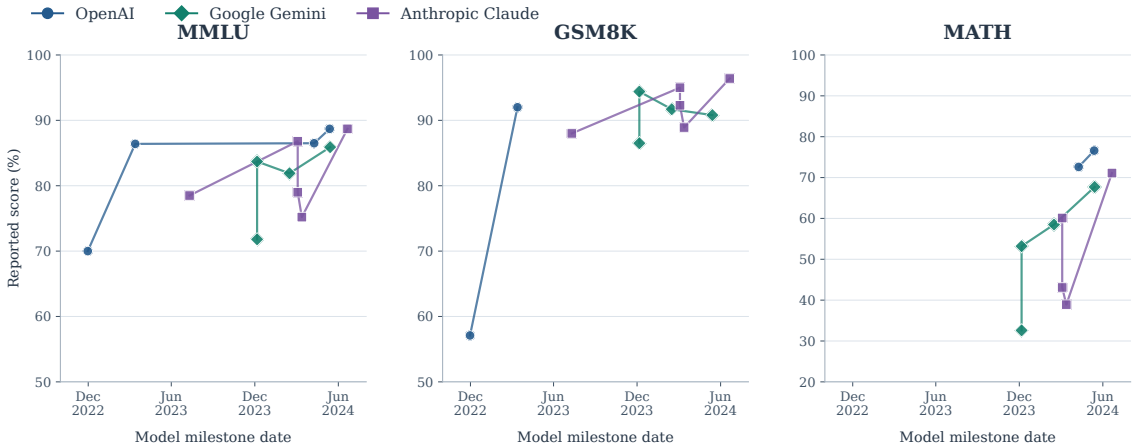


OpenAI: GPT-3.5 → GPT-4 → GPT-4 Turbo → GPT-4o

Google Gemini: 1.0 Pro / Ultra → 1.5 Pro (February preview → May update)

Anthropic Claude: 2 → 3 Opus / Sonnet → 3 Haiku → 3.5 Sonnet

# Model performance over time: scientific and multidisciplinary reasoning

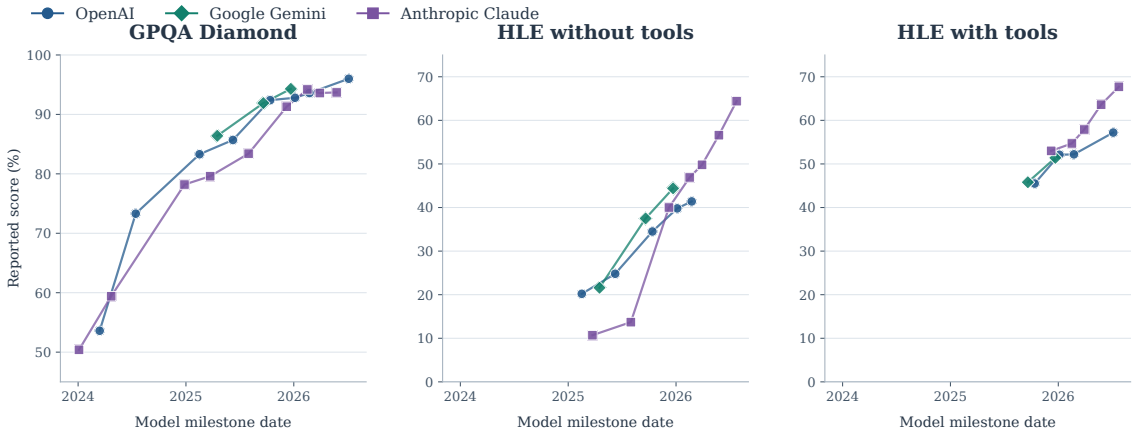


OpenAI: GPT-4o → o1-preview → o3 → GPT-5 → 5.2 → 5.4 → 5.5 → GPT-6 Astra

Google Gemini: 2.5 Pro → 3 Pro → 3.1 Pro

Anthropic Claude: 3 Opus → 3.5 Sonnet → 3.7 Sonnet → Opus 4 → Sonnet 4.5 → Opus 4.6 → 4.7 → 4.8 → 5 → 5.5

## From academic questions to complete workflows

### 2019

- **Nov - ARC-AGI-1**
  Learning

### 2020

- **Sep - MMLU**
  Knowledge

### 2021

- **Mar - MATH**
  Mathematics
- **Jul - HumanEval**
  Software
- **Oct - GSM8K**
  Mathematics

### 2023

- **Jul - WebArena**
  Web / desktop
- **Oct - SWE-bench**
  Software
- **Nov - GPQA / Diamond**
  Knowledge
- **Nov - MMMU**
  Multimodal

### 2024

- **Apr - OSWorld**
  Web / desktop
- **Jun - MMLU-Pro**
  Knowledge
- **Jul - LAB-Bench**
  Science
- **Aug - SWE-bench Verified**
  Software
- **Sep - MMMU-Pro**
  Multimodal
- **Nov - FrontierMath**
  Mathematics

### 2025

- **Jan - Humanity's Last Exam**
  Knowledge
- **Mar - ARC-AGI-2**
  Learning
- **Apr - BrowseComp**
  Web / desktop
- **May - Terminal-Bench**
  Terminal
- **Sep - GDPval**
  Professional work
- **Nov - Terminal-Bench 2.0**
  Terminal

## The widening benchmark landscape in 2026

### Jan-Mar

- **Jan - FrontierMath Open Problems**
  Mathematics
- **Feb - LABBench2**
  Science
- **Mar - Riemann-bench**
  Mathematics
- **Mar - ARC-AGI-3**
  Learning

### Apr-Jun

- **Apr - AutomationBench**
  Professional work
- **May - DeepSWE v1**
  Software
- **Jun - Agents' Last Exam**
  Professional work
- **Jun - FrontierMath Tier 4 v2**
  Mathematics
- **Jun - DeepSWE v1.1**
  Software
- **Jun - OSWorld 2.0**
  Web / desktop

### Jul-Sep

- **Jul - AutomationBench v1.0.6**
  Professional work
- **Aug - Terminal-Bench-Science 0.1**
  Science
- **Aug - Terminal-Bench 4.0**
  Terminal
- **Sep - FrontierSWE v2**
  Software
- **Sep - OSWorld 2.1**
  Web / desktop
- **Sep - HLE-Diamond**
  Knowledge

### October

- **Oct - Agents' Last Exam v1.1**
  Professional work
- **Oct - FrontierCode 2.0**
  Software
- **Oct - Artificial Analysis v5.0**
  Benchmark adoption

# Increasing capability sharpens the judgement question

| Reported benchmark score (%) | Earlier milestone | Later milestone |
|---|---|---|
| MMLU: GPT-3.5 to GPT-4o | 70.0 | **88.7** |
| GPQA Diamond: GPT-4o to GPT-6 Astra | 53.6 | **96.0** |

These tests reward answers against a defined key. Accounting estimates and decisions also require a criterion for how evidence should be weighted.

Which criterion matters: the facts, an accounting rule, economic consistency, or the intended user's response?

# Koriat (2008): Answer switching

- Koriat (2008): 41 students answered 105 two-choice general-knowledge questions twice.
- After a short break, the same questions appeared in a new order.
- Participants repeated **83.4%** of their initial answers.

The implied average switching rate is **16.6%**.

# Koriat (2012): Confidence and switching

| Initial confidence | Repeat answer | Switch answer |
|---|---|---|
| Below own median | 74% | **26%** |
| Above own median | 95% | **5%** |

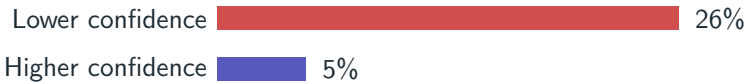

# Goldberg (1970): Modelling judgement

- Goldberg (1970) analyzed 29 clinical psychologists judging 861 MMPI profiles.
- For each psychologist, a linear model predicted that psychologist's rating from 11 MMPI scales.
- The fitted coefficients used the psychologist's judgments as the target.

# Goldberg (1970): Criterion accuracy

| Predictor | Correlation with recorded diagnosis |
|---|---|
| Psychologists | .28 |
| Models of psychologists | **.31** |

In split-sample cross-validation, models outperformed their corresponding psychologists in **86%** of comparisons.

# Benchmarks for evaluating a judgment

| Outcome | Benchmark |
|---|---|
| Factual accuracy | The stated amount, classification, or calculation |
| Rule compliance | The governing accounting requirement |
| Economic consistency | A specified payoff, ordering, or economic criterion |
| Human recovery | The response of the intended person or group |

# OpenAI interaction style and behavioral changes

# Useful information requires accounting judgement

| Framework characteristic | Requirement |
|---|---|
| Relevance | Predictive or confirmatory value for decisions. |
| Faithful representation | Complete, neutral, and free from error. |
| Enhancing qualities | Comparability, verifiability, timeliness, understandability. |
| Cost constraint | Reporting benefits justify the costs. |

Estimates incorporate future expectations and estimation error (Barth, 2006; Dechow and Dichev, 2002). Evaluating complex estimates requires integrating evidence (Griffith et al., 2015).

# The development of benchmarks for tabular data

**FinQA**

Financial tables plus numerical programs

**TAT-QA**

Tables and narrative disclosures

**TableBench preprint**

General table reasoning and visualization

**TableBench at AAAI**

A broad table-reasoning benchmark

**ToRR**

Robustness across representations

**TabArena**

Prediction from structured features

**ToRR revision**

Updated representation tests

**SpreadsheetBench 2**

321 professional workbook tasks

# Atomicity of tabular information

Kim, Kang, and Jung (2026) decompose table structure into six atomic features.

| Atomic feature | Relationship conveyed |
|---|---|
| Grid formation | Which value occupies each row–column intersection. |
| Header designation | Which cells describe the data cells. |
| Cell merging | Which labels span rows or columns. |
| Header nesting | How categories form a hierarchy. |
| Section stacking | How sections share column structures. |
| Table composition | How different table structures fit together. |

# The expansion of context-window capacity

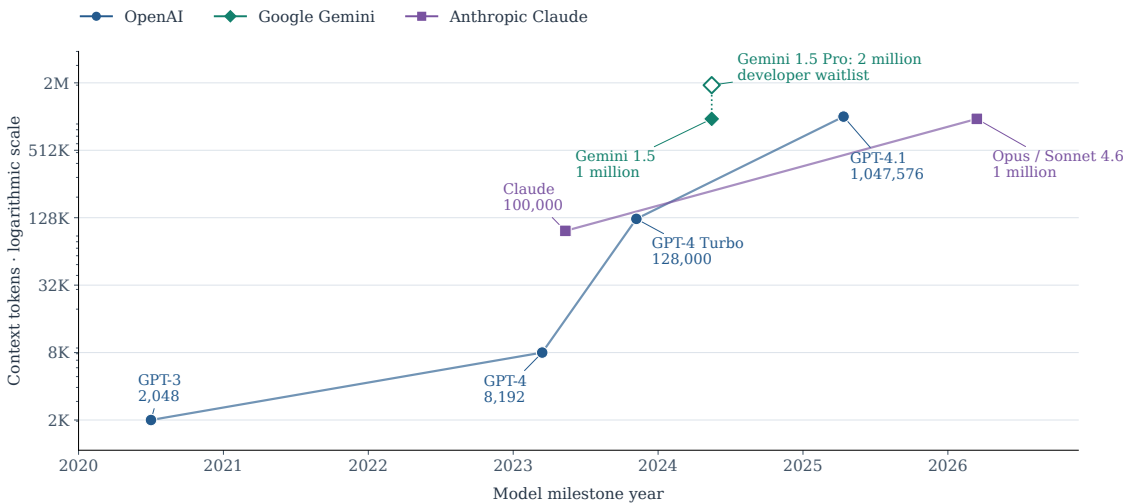

# Users bring knowledge and preferences to the same case

- Knowledge affects which classifications, relationships, and disclosures a user recognizes.
- Preferences affect how the user weighs uncertainty and decision consequences.
- A participant group combines people with different characteristics.

# Predictions for accounting judgements

**H1: What does the reporting choice change?**

Within presentation-based experiments, AI recovers human responses more often when the treatment changes information than when it changes presentation while holding information constant.

**H2: Is correct processing sufficient?**

AI and human judgements can differ even when AI correctly answers separate questions about the relevant accounting information.

# Hold the accounting case fixed; describe the decision maker

- Fix the source-defined role, treatments, questions, and response scale.
- Vary assigned knowledge, preferences, and other participant characteristics.
- Compare the same finding under neutral and profile instructions.

# The source catalogue and the analytical samples

| Analytical population | Count | Papers |
|---|---|---|
| Source catalogue: 401 accepted records | 395 | 64 |
| Source-compatible nonzero directions | 208 | 47 |
| Controlled procedure patterns | 45 | 28 |
| Unprofiled model-comparison panel | 236 | 46 |
| Findings with complete human mean benchmarks | 84 | 23 |

The inventory contains 65 studies. Paired tests use the same available findings across the compared profiles.

# Participant profiles describe the decision maker

| Profile family | Profiles | Description |
|---|---|---|
| Reference prompts | 2 | Original prompt; fresh neutral prompt |
| Assigned knowledge | 1 | Accounting expertise |
| Uncertainty and evidence | 8 | Cautious/open attitudes; higher/ordinary thresholds; with or without knowledge |
| Other characteristics | 8 | Horizons, familiarity, attention, objectives |
| Financial risk | 4 | SCF investment-risk categories |

Primary profile model: GPT-6 Luna; temperature 0; reasoning disabled.

# Four risk profiles define the audience mixture

Substantial risk; above-average risk; average risk; no financial risk.

$$\bar{y}_g(t) = \sum_{j=1}^{4} w_{gj} y_j(t), \qquad \sum_{j=1}^{4} w_{gj} = 1.$$

| Audience mixture | Shares applied to the same four responses |
|---|---|
| Equal mixture | 25% in each risk category |
| Dated audience | Closest preceding Survey of Consumer Finances wave |
| Fixed audience | 2022 Survey of Consumer Finances distribution |

# Directions, judgement levels, and economic decisions

| Outcome | Evaluation criterion |
|---|---|
| Human recovery | Same sign as the source-compatible nonzero human contrast |
| Controlled pattern agreement | Source-defined pattern under the adapted task procedure |
| Judgement-level agreement | Distance from published human condition means |
| Economic agreement | Source-supported accounting rule or economic decision |

Separate factual checks evaluate amounts, classifications, and relationships.

# Two additional accounting studies test transfer

| Study | Initial decision task | Conditions |
|---|---|---|
| Le (2026) | Accounting rules and accounting-career judgements | 3 |
| Cai et al. (2026) | Cryptocurrency classification, measurement, and valuation | 2 |

Apply the 23 participant profiles to the five initial decision screens.

These applications use the profile descriptions without audience weights.

# The empirical sequence

1. Establish recovery on the accounting tasks.
2. Vary participant knowledge and characteristics.
3. Combine four risk-profile responses and vary their audience shares.
4. Assess judgement levels and the two additional accounting studies.
5. Examine robustness of the accounting comparisons.

# References i

Lisa P. Argyle, Ethan C. Busby, Nancy Fulda, Joshua R. Gubler, Christopher Rytting, and David Wingate. Out of One, Many: Using Language Models to Simulate Human Samples. *Political Analysis*, 31(3):337–351, July 2023. ISSN 1047-1987, 1476-4989. doi: 10.1017/pan.2023.2. URL https://www.cambridge.org/core/journals/political-analysis/article/out-of-one-many-using-language-models-to-simulate-human-samples/035D7C8A55B237942FB6DBAD7CAA4E49.

Mary E. Barth. Including estimates of the future in today's financial statements. *Accounting Horizons*, 20(3):271–285, 2006. doi: 10.2308/acch.2006.20.3.271.

Sarah E. Bonner and Barry L. Lewis. Determinants of auditor expertise. *Journal of Accounting Research*, 28(Supplement):1–20, 1990. doi: 10.2307/2491243. URL https://doi.org/10.2307/2491243.

Eddy Cardinaels. The interplay between cost accounting knowledge and presentation formats in cost-based decision-making. *Accounting, Organizations and Society*, 33(6):582–602, 2008. doi: 10.1016/j.aos.2007.06.003. URL https://doi.org/10.1016/j.aos.2007.06.003.

# References ii

C. Bryan Cloyd. The effects of financial accounting conformity on recommendations of tax preparers. *The Journal of the American Taxation Association*, 17(2):50, 1995. ISSN 01989073. URL https://www.proquest.com/docview/211035744/abstract/86182E79C5454ABCPQ/1. Num Pages: 50 Place: Sarasota, United States Publisher: American Accounting Association.

Thomas R. Cook, Sophia Kazinnik, Zach Modig, and Nathan M. Palmer. What do LLMs want? Finance and Economics Discussion Series 2026-006, Board of Governors of the Federal Reserve System, 2026. URL https://www.federalreserve.gov/econres/feds/what-do-llms-want.htm.

Patricia M. Dechow and Ilia D. Dichev. The quality of accruals and earnings: The role of accrual estimation errors. *The Accounting Review*, 77(Supplement):35–59, 2002. doi: 10.2308/accr.2002.77.s-1.35.

Danica Dillion, Niket Tandon, Yuling Gu, and Kurt Gray. Can ai language models replace human participants? *Trends in Cognitive Sciences*, 27(7):597–600, 2023. doi: 10.1016/j.tics.2023.04.008.

Wenyu Du, Rui Zheng, Tongxu Luo, Stephen Chung, and Jie Fu. Fine-grained manipulation of arithmetic neurons. In *Proceedings of the 8th BlackboxNLP Workshop: Analyzing and Interpreting Neural Networks for NLP*, pages 467–479. Association for Computational Linguistics, 2025. URL https://aclanthology.org/2025.blackboxnlp-1.27/.

# References iii

Carsten Erb and Christoph Pelger. "Twisting words"? A study of the construction and reconstruction of reliability in financial reporting standard-setting. *Accounting, Organizations and Society*, 40:13–40, January 2015. ISSN 0361-3682. doi: 10.1016/j.aos.2014.11.001. URL https://www.sciencedirect.com/science/article/pii/S0361368214001007.

Marc Eulerich, Aida Sanatizadeh, Hamid Vakilzadeh, and David A. Wood. Is it all hype? chatgpt's performance and disruptive potential in the accounting and auditing industries. *Review of Accounting Studies*, 29(3):2318–2349, 2024. doi: 10.1007/s11142-024-09833-9.

FASB. Conceptual framework for financial reporting. Technical report, Financial Accounting Standards Board, 2024. URL https://storage.fasb.org/Conceptual%20Framework%20for%20Financial%20Reporting%20%28September%202024%29.pdf. September 2024 compilation, Chapters 1–8.

James R. Frederickson and Jeffrey S. Miller. The effects of pro forma earnings disclosures on analysts' and nonprofessional investors' equity valuation judgments. *The Accounting Review*, 79(3):667–686, 2004. doi: 10.2308/accr.2004.79.3.667. URL https://doi.org/10.2308/accr.2004.79.3.667.

Oscar Gilg, Pierre Beckmann, Daniel Paleka, and Patrick Butlin. Probing persona-dependent preferences in language models, 2026. URL https://arxiv.org/abs/2605.13339. arXiv:2605.13339.

# References iv

Lewis R. Goldberg. Man versus model of man: A rationale, plus some evidence, for a method of improving on clinical inferences. *Psychological Bulletin*, 73(6):422–432, 1970. ISSN 1939-1455. doi: 10.1037/h0029230.

E. E. Griffith, J. S. Hammersley, K. Kadous, and D. Young. Auditor mindsets and audits of complex estimates. *Journal of Accounting Research*, 53(1):49–77, 2015. doi: 10.1111/1475-679X.12066.

D. Eric Hirst, Patrick E. Hopkins, and James M. Wahlen. Fair Values, Income Measurement, and Bank Analysts' Risk and Valuation Judgments. *The Accounting Review*, 79(2):453–472, 2004. ISSN 0001-4826. URL https://www.jstor.org/stable/3203252. Publisher: American Accounting Association.

Chaemin Jang, Dongman Lee, and Jihee Kim. Instruction-tuned language models cannot sample from distributions they can describe, 2026. URL https://arxiv.org/abs/2607.25292. arXiv:2607.25292.

# References v

Hang Jiang, Xiajie Zhang, Xubo Cao, Cynthia Breazeal, Deb Roy, and Jad Kabbara. PersonaLLM: Investigating the ability of large language models to express personality traits. In *Findings of the Association for Computational Linguistics: NAACL 2024*, pages 3605–3627, 2024. doi: 10.18653/v1/2024.findings-naacl.229. URL https://aclanthology.org/2024.findings-naacl.229/.

Eunbin Kim, Hyeonseok Kang, and Sangkeun Jung. LTEM: LLM-friendly table-enhanced markdown for table understanding and generation. In *Proceedings of the 2026 Conference on Empirical Methods in Natural Language Processing*, pages 43304–43324, 2026. doi: 10.18653/2026.emnlp-main.1900. URL https://aclanthology.org/2026.emnlp-main.1900/.

Lisa Koonce, Marlys Gascho Lipe, and Mary Lea McAnally. Judging the Risk of Financial Instruments: Problems and Potential Remedies. *The Accounting Review*, 80(3):871–895, 2005. ISSN 0001-4826. URL https://www.jstor.org/stable/4093180. Publisher: American Accounting Association.

Lisa Koonce, Cassie Mongold, Laura Quaid, and Brian J. White. Experimental research on standard-setting issues in financial reporting. *Accounting, Organizations and Society*, 112:101509, June 2024. ISSN 0361-3682. doi: 10.1016/j.aos.2023.101509. URL https://www.sciencedirect.com/science/article/pii/S0361368223000806.

# References vi

Asher Koriat. Subjective confidence in one's answers: The consensuality principle. *Journal of Experimental Psychology: Learning, Memory, and Cognition*, 34(4):945–959, 2008. doi: 10.1037/0278-7393.34.4.945.

Asher Koriat. The self-consistency model of subjective confidence. *Psychological Review*, 119(1): 80–113, 2012. doi: 10.1037/a0025648.

Bradford Levy. Caution ahead: Numerical reasoning and look-ahead bias in AI models. *Journal of Accounting Research*, 64(3):1139–1188, 2026. doi: 10.1111/1475-679x.70058. URL https://doi.org/10.1111/1475-679x.70058.

Robert Libby and Joan Luft. Determinants of judgment performance in accounting settings: Ability, knowledge, motivation, and environment. *Accounting, Organizations and Society*, 18(5):425–450, 1993. doi: 10.1016/0361-3682(93)90040-D. URL https://doi.org/10.1016/0361-3682(93)90040-D.

Nikki L. MacKenzie, James Moon, Susan Rykowski, Quinn Thomas Swanquist, and Robert Lowell Whited. A Comparison of Artificial Intelligence and Human Responses in Audit Experiments, July 2026. URL https://papers.ssrn.com/abstract=5036687. Working paper, revised July 8, 2026; first posted January 6, 2025.

# References vii

Paul B. W. Miller and David Mosso. Financial accounting measurement: Why things may not be what they seem to be. *Financial Executive*, pages 28–35, December 1983.

Christoph Pelger. Practices of standard-setting – An analysis of the IASB's and FASB's process of identifying the objective of financial reporting. *Accounting, Organizations and Society*, 50:51–73, April 2016. ISSN 0361-3682. doi: 10.1016/j.aos.2015.10.001. URL https://www.sciencedirect.com/science/article/pii/S0361368215000896.

Shibani Santurkar, Esin Durmus, Faisal Ladhak, Cinoo Lee, Percy Liang, and Tatsunori Hashimoto. Whose opinions do language models reflect? In *Proceedings of the 40th International Conference on Machine Learning*, volume 202 of *Proceedings of Machine Learning Research*, pages 29971–30004, 2023. URL https://proceedings.mlr.press/v202/santurkar23a.html.